

Google Cloud Tutorial

BEFORE WE BEGIN

BIG REMINDER: Make sure you stop your instances!

(We know you won't read until the very bottom once your assignment is running, so we are printing this at the top too since it is ***super important***)

Don't forget to ***stop your instance*** when you are done (by clicking on the stop button at the top of the page showing your instances), otherwise you will ***run out of credits*** and that will be very sad. :(

If you follow our instructions below correctly, you should be able to restart your instance and the downloaded software will still be available.

Left my instance running for no reason



No more Google Cloud credits

[Create and Configure Your Account](#)

For the class project and assignments, we offer an option to use Google Compute Engine for developing and testing your implementations. This tutorial lists the necessary steps of working on the assignments using Google Cloud. **We expect this tutorial to take about an hour. Don't get intimidated by the steps, we tried to make the tutorial detailed so that you are less likely to get stuck on a particular step. Please tag all questions related to Google Cloud with google_cloud on Piazza.**

This tutorial goes through how to set up your own Google Compute Engine (GCE) instance to work on the assignments. Each student will have

100 *in credit throughout the quarter*. When you sign up for the first time, you also receive 300 credits from Google by default. Please try to use the resources judiciously. But if \$100 ends up not being enough, we will try to adjust this number as the quarter goes on.

First, if you don't have a Google Cloud account already, create one by going to the [Google Cloud homepage](#) and clicking on **Compute**. When you get to the next page, click on the blue **TRY IT FREE** button. If you are not logged into gmail, you will see a page that looks like the one below. Sign into your gmail account or create a new one if you do not already have an account.

The screenshot shows the initial landing page for trying Google Cloud Platform for free. It features a blue header bar with the text "Try Cloud Platform for free" and the Google logo. Below the header, there's a "Country" dropdown menu set to "United States". Under the "Acceptances" section, there's a checkbox for receiving updates and two radio buttons for "Yes" and "No". A note states: "I agree that my use of any services and related APIs is subject to my compliance with the applicable Terms of Service. I have also read and agree to the Google Cloud Platform Free Trial Terms of Service." Below this, another radio button for "Yes" is selected. At the bottom of the main form is a blue "Agree and continue" button. To the right of the main form, there's a sidebar with the heading "You won't be charged" and a note: "Your payment information is used to verify your identity until you upgrade". The background of the page has a hexagonal grid pattern.

Click the appropriate **yes** or **no** button for the first option, and check **yes** for the second option after you have read the required agreements. Press the blue **Agree and continue** button to continue to the next page to enter the requested information (your name, billing address and credit card information). Remember to select "**Individual**" as "Account Type":

Try Cloud Platform for free

Google

Customer info

Account type Individual

Name and address

Name: Tony Stark

Address line 1: 123 Avenger Avenue

Address line 2:

City: Stanford

State: California ZIP code: 94305

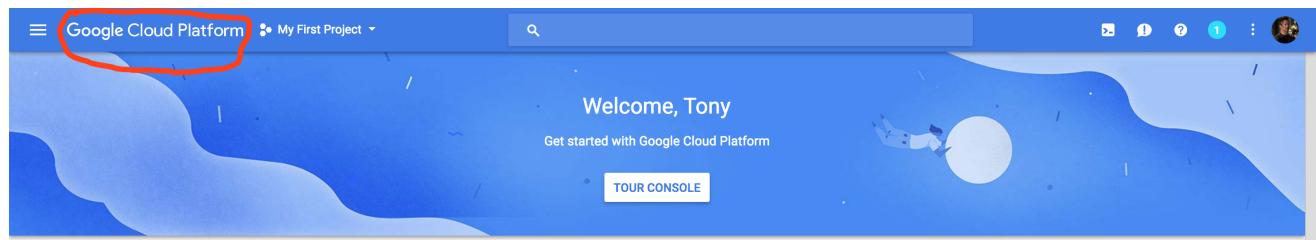
Phone number: 4082223333

You won't be charged

Your payment information is used to verify your identity until you upgrade



Once you have entered the required information, press the blue **Start my free trial** button. You will be greeted by a page like this:



Welcome, Tony

Get started with Google Cloud Platform

[TOUR CONSOLE](#)

Explore

- Cloud Launcher**: Explore, launch, and manage solutions in just a few clicks.
- API**: Google APIs. Enable APIs, create credentials, and track your usage.
- Documentation**: Explore Google Cloud Platform documentation.
- Create an empty project**

Quick starts

- Try Compute Engine**: Create a Linux virtual machine instance in Compute Engine in this guided walkthrough.
- Build a guestbook on Kubernetes Engine**: Deploy a Guestbook application with Google Kubernetes Engine and Kubernetes.
- Try App Engine**: Create and deploy a Hello World app using the Go programming language.
- Learn to use Cloud Storage**: Create a storage bucket and add a publicly accessible file.

[SHOW ALL](#)

Press the "Google Cloud Platform" (in red circle), and it will take you to the main dashboard:

The screenshot shows the Google Cloud Platform dashboard for a project named "My First Project". The left sidebar contains links for Home, Cloud Launcher, Billing, APIs & Services, Support, IAM & admin, Getting started, Compute Engine, App Engine, Compute Engine, Kubernetes Engine, Cloud Functions, Storage, Bigtable, and Datastore. The main content area is divided into several sections: "Project info" (Project name: My First Project, Project ID: analog-stage-200103, Project number: 62970935707), "APIs" (Shows requests per second from 7:30 to 8:15), "Resources" (Shows no resources), "Trace" (Shows no trace data from the past 7 days), "Getting Started" (Shows how to enable APIs and get credentials like keys, and a link to Deploy a prebuilt solution), and "Google Cloud Platform status" (Shows all services normal). A "Billing" section indicates estimated charges of \$0.00 for the period April 1 - 3, 2018. An "Error Reporting" section shows no errors. A "News" section has a link to learn how to set up Error Reporting.

To change the name of your project, click on [Go to project settings](#) under the **Project info** section.

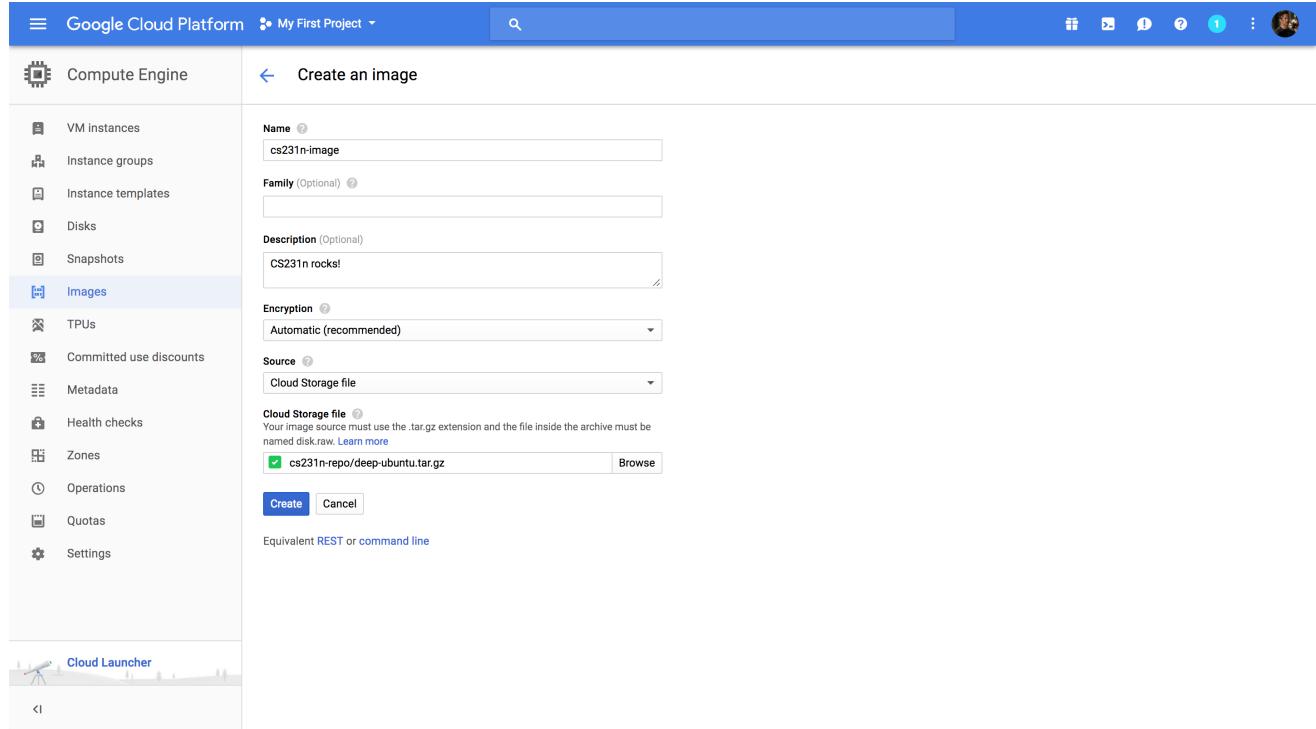
Create an image from our provided disk

For all assignments and the final project, we provide you with a pre-configured disk that contains the necessary environment and deep learning frameworks. To use our disk, you first need to create your own custom image using our file, and use this custom image as the boot disk for your new VM instance.

Go to **Compute Engine**, then **Images** and click on the blue **Create Image** button at the top of the page. See the screenshot below.

The screenshot shows the "Compute Engine" section of the Google Cloud Platform interface, specifically the "Images" page. The left sidebar includes links for VM instances, Instance groups, Instance templates, Disks, Snapshots, Images (which is selected and highlighted in blue), TPUs, Committed use discounts, Metadata, Health checks, Zones, Operations, Quotas, and Settings. The main content area features a table of existing images and a "CREATE IMAGE" button. The table columns are Name, Size, Created by, and Family. The images listed include cs231n, cs231n-old, centos-6-v20180401, centos-7-v20180401, coreos-alpha-1729-0-0-v20180327, coreos-beta-1722-2-0-v20180330, coreos-stable-1688-5-3-v20180403, cos-beta-66-10452-28-0, cos-dev-67-10527-0-0, cos-stable-64-10176-62-0, cos-stable-65-10323-64-0, debian-8-jessie-v20180401, debian-9-stretch-v20180401, rhel-6-v20180401, rhel-7-sap-apps-v20180401, rhel-7-sap-hana-v20180401, and rhel-7-v20180401. To the right of the table is a panel titled "Select an image" with a "LABELS" section indicating "No images selected".

Enter your preferred name in the **Name** field. Mine is called **cs231n-image**. Select **Cloud Storage file** for **Source**, enter **cs231n-repo/deep-ubuntu.tar.gz** and click on the blue **Create** button. See the screenshot below. It will take a few minutes for your image to be created (about 10-15 in our experience, though your mileage may vary).



Launch a Virtual Instance

To launch a virtual instance, go to the **Compute Engine** menu on the left column of your dashboard and click on **VM instances**.

Then click on the blue **Create** button on the next page. This will take you to a page that looks like the screenshot below. **(NOTE: Please carefully read the instructions in addition to looking at the screenshots. The instructions tell you exactly what values to fill in).**

Make sure that the Zone is set to be **us-west1-b** (especially for assignments where you need to use GPU instances). Under **Machine type** pick the **8 vCPUs** option. Click on the **customize** button under **Machine type** and make sure that the number of cores is set to 8 and the number of GPUs is set to **None** (we will not be using GPUs in assignment 1. GPU will be covered later in this tutorial).

Click on the **Change** button under **Boot disk**, choose **Custom images**, you will see this screen:

Select the image you created in the previous step, here it's **cs231n-image**. Also increase the boot disk size as you see fit. Click **Select** and you will get back to the "create instance" screen.

Check **Allow HTTP traffic** and **Allow HTTPS traffic**. Expand the **Management, disks, networking, SSH keys** menu if it isn't visible, select **Disks** tab, and uncheck **Delete boot disk when instance is deleted**.

The screenshot shows the 'Create an instance' wizard in the Google Cloud Platform. On the left, a sidebar lists various Compute Engine options like VM instances, Instance groups, and Disks. The main form is for creating a new VM instance. Under the 'Disks' tab, a 'Boot disk' section shows a 'New 40 GB standard persistent disk' named 'Image cs231n-image'. Below this, the 'Identity and API access' section includes a 'Service account' dropdown set to 'Compute Engine default service account' and 'Access scopes' set to 'Allow default access'. The 'Firewall' section has two checked boxes: 'Allow HTTP traffic' and 'Allow HTTPS traffic'. At the bottom of the disk configuration, there's a 'Deletion rule' section with a checkbox labeled 'Delete boot disk when instance is deleted', which is circled in red. Other tabs like Management, Networking, and SSH Keys are also visible.

Click on the blue **Create** button at the bottom of the page. You should have now successfully created a Google Compute Instance, it might take a few minutes to start running. When the instance is ready, your screen should look something like the one below. When you want to stop running the instance, click on the blue stop button above.

The screenshot shows the 'VM instances' page in the Google Cloud Platform. The sidebar on the left is identical to the previous screenshot. The main area displays a table of VM instances. The first row, for 'cs231', is selected and shows its details: Zone 'us-west1-b', Internal IP '10.138.0.2', External IP '35.185.240.182', and a 'Connect' column with 'SSH' and a more options menu. There is also a 'Columns' dropdown menu. A 'CREATE INSTANCE' button is visible at the top of the list.

Take note of your instance name, you will need it to ssh from your laptop.

Connect to Your Virtual Instance

Now that you have created your virtual GCE, you want to be able to connect to it from your computer. The rest of this tutorial goes over how to do that using the command line. First, download the Google Cloud SDK that is appropriate for your platform from [here](#) and follow their installation instructions. **NOTE: this tutorial assumes that you have performed step #4 on the website which they list as optional.** When prompted, make sure you select `us-west1-b` as the time zone.

The easiest way to connect is using the gcloud compute command below. The tool takes care of authentication for you. On your laptop (OS X for example), run:

```
gcloud compute ssh --zone=us-west1-b <YOUR-INSTANCE-NAME>
```

If `gcloud` command is not in system path, you can also reference it by its full path `/<DIRECTORY-WHERE-GOOGLE-CLOUD-IS-INSTALLED>/bin/gcloud`. See [this page](#) for more detailed instructions.

First time setup

Upon your first ssh, you need to run a one-time setup script and reload the `.bashrc` to activate the libraries. The exact command is

```
/home/shared/setup.sh && source ~/.bashrc
```

The command will download a git repo, patch your `.bashrc` and copy a jupyter notebook config file to your home directory. If you ever switch account/username, you will have to re-run the setup command. If you see any permission error, simply prepend `sudo` to the command.

When the command finishes without error, run `which python` on the command line and it should report `/home/shared/anaconda3/bin/python`. See screenshot:

```
tonystark@cs231:~$ which python
/home/shared/anaconda3/bin/python
tonystark@cs231:~$ python
Python 3.6.4 |Anaconda custom (64-bit)| (default, Jan 16 2018, 18:10:19)
[GCC 7.2.0] on linux
Type "help", "copyright", "credits" or "license" for more information.
>>> import tensorflow
/home/shared/anaconda3/lib/python3.6/site-packages/h5py/_init__.py:36: FutureWarning: Conversion of the second argument of issubdtype from `float` to `np.floating` is deprecated. In future, it will be treated as `np.float
64 == np.dtype(float).type`.
  from ._conv import register_converters as _register_converters
>>> import torch
>>> import keras
Using TensorFlow backend.
>>> import caffe2
>>> █
```

(don't worry about the Tensorflow warning message)

Our provided image supports the following frameworks:

- [Anaconda3](#), a python package manager. You can think of it as a better alternative to `pip`.
- Numpy, matplotlib, and tons of other common scientific computing packages.
- [Tensorflow 1.7](#), both CPU and GPU.
- [PyTorch 0.3](#), both CPU and GPU.
- [Keras](#) that works with Tensorflow 1.7
- [Caffe2](#), CPU only. Note that it is very different from the original Caffe.

- Nvidia runtime: CUDA 9.0 and cuDNN 7.0. They only work when you create a Cloud GPU instance, which we will cover later.

The `python` on our image is `3.6.4`, and has all the above libraries installed. It should work out of the box for all assignments unless noted otherwise. You don't need `virtualenv`, but if you insist, Anaconda has [its own way](#). If you need libraries not mentioned above, you can always run `conda install <mylib>` yourself.

You are now ready to work on the assignments on Google Cloud!

Using Jupyter Notebook with Google Compute Engine

Many of the assignments will involve using Jupyter Notebook. Below, we discuss how to run Jupyter Notebook from your GCE instance and connect to it with your local browser.

Getting a Static IP Address

Change the External IP address of your GCE instance to be static (see screenshot below).

Name	Zone	Recommendation	Internal IP	External IP	Connect
instance-1	us-west1-b		10.138.0.2	35.185.223.167	SSH

To Do this, click on the 3 line icon next to the **Google Cloud Platform** button on the top left corner of your screen, go to **VPC network** and **External IP addresses** (see screenshot below).

VM instances

CREATE INSTANCE IMPORT VM REFRESH START STOP RESET SHOW INFO PANEL

Name	Zone	Recommendation	Internal IP	External IP	Connect
cs231	us-west1-b		10.138.0.2	35.185.240.182	SSH

Filter VM instances Columns

Storage

- Bigtable
- Datastore
- Storage
- SQL
- Spanner

Networking

- VPC network
 - VPC networks
 - External IP addresses
 - Firewall rules
 - Routes
 - VPC network peering
 - Shared VPC
- Network services
- Hybrid Connectivity
- Network Security

StackDriver

- Monitoring
- Debug
- Trace
- Logs

To have a static IP address, change **Type** from **Ephemeral** to **Static**. Enter your preferred name for your static IP, ours is `cs231n-ip` (see screenshot below). And click on Reserve. Remember to release the static IP address when you are done because according to [this page](#) Google charges a small fee for unused static IPs.

External IP addresses

+ RESERVE STATIC ADDRESS REFRESH RELEASE STATIC ADDRESS SHOW INFO PANEL

Name	External Address	Region	Type	Version	In use by	Labels
-	35.185.240.182	us-west1	Static	IPv4	VM instance cs231 (Zone b)	Change

Reserve a new static IP address

Name: cs231n-ip

Description (Optional):

CANCEL RESERVE

VPC network

External IP addresses

Firewall rules

Routes

VPC network peering

Shared VPC

Take note of your Static IP address (circled on the screenshot below). We use 35.185.240.182 for this tutorial.

The screenshot shows the Google Cloud Platform interface for managing VPC networks. On the left, there's a sidebar with options like 'VPC networks', 'External IP addresses' (which is selected and highlighted in blue), 'Firewall rules', 'Routes', 'VPC network peering', and 'Shared VPC'. The main content area is titled 'External IP addresses' and shows a table with one row. The columns are 'Name', 'External Address' (with a red circle around it), 'Region', 'Type', 'Version', 'In use by', and 'Labels'. The single row contains 'cs231n-ip' and '35.185.240.182' respectively. There are buttons for 'RESERVE STATIC ADDRESS', 'REFRESH', and 'RELEASE STATIC ADDRESS' at the top, and a 'SHOW INFO PANEL' button on the right.

Adding a Firewall rule

One last thing you have to do is adding a new firewall rule allowing TCP access to a particular port number. The default port we use for Jupyter is **7000**. You can find this default value in the config file generated at setup time (`~/.jupyter/jupyter_notebook_config.py`). Feel free to change it.

Click on the 3-line icon at the top of the page next to **Google Cloud Platform**. On the menu that pops up on the left column, go to **VPC network** and **Firewall rules** (see the screenshot below).

The screenshot shows the Google Cloud Platform interface for managing Firewall rules. The left sidebar has the same navigation as the previous screenshot. The main content area is titled 'Firewall rules' and shows a table of rules. The columns are 'Name', 'Targets', 'Source filters', 'Protocols / ports', 'Action', 'Priority', and 'Network'. There are buttons for '+ CREATE FIREWALL RULE', 'REFRESH', and 'DELETE' at the top. The table lists several default rules: 'default-allow-http' (http-server, IP ranges: 0.0.0.0/0, tcp:80, Allow, Priority 1000, Network default), 'default-allow-https' (https-server, IP ranges: 0.0.0.0/0, tcp:443, Allow, Priority 1000, Network default), 'default-allow-icmp' (Apply to all, IP ranges: 0.0.0.0/0, icmp, Allow, Priority 65534, Network default), 'default-allow-internal' (Apply to all, IP ranges: 10.128.0.0/9, tcp:0-65535, udp:0-65535, 1 more, Allow, Priority 65534, Network default), 'default-allow-rdp' (Apply to all, IP ranges: 0.0.0.0/0, tcp:3389, Allow, Priority 65534, Network default), and 'default-allow-ssh' (Apply to all, IP ranges: 0.0.0.0/0, tcp:22, Allow, Priority 65534, Network default). A note in the sidebar says: 'Firewall rules control incoming or outgoing traffic to an instance. By default, incoming traffic from outside your network is blocked. [Learn more](#)'.

Click on the blue **CREATE FIREWALL RULE** button. Enter whatever name you want: we use cs231n-rule. Select "All instances in the network" for **Targets** (if the menu item exists). Enter `0.0.0.0/0` for **Source IP ranges** and `tcp:<port-number>` for **Specified protocols and ports** where `<port-number>` is the number you used above. Click on the blue **Create** button. See the screenshot below.

Launching and connecting to Jupyter Notebook

After you ssh into your GCE instance using the prior instructions, run Jupyter notebook from the folder with your assignment files. As a quick example, let's launch it from `/home/shared` folder.

```
cd /home/shared
jupyter-notebook --no-browser --port=7000
```

If you simply run `jupyter-notebook` without any command line arguments, it will pick up the default config values in `~/.jupyter/jupyter_notebook_config.py`. In our disk image, it is `no-browser` and port 7000 by default.

The command should block your stdin and display something like:

```
tonystark@cs231:~/home/shared$ jupyter-notebook --no-browser --port=7000
[I 01:17:20.105 NotebookApp] Writing notebook server cookie secret to /run/user/1015/jupyter/notebook_cookie_secret
[W 01:17:20.519 NotebookApp] WARNING: The notebook server is listening on all IP addresses and not using encryption. This is not recommended.
[I 01:17:20.550 NotebookApp] JupyterLab beta preview extension loaded from /home/shared/anaconda3/lib/python3.6/site-packages/jupyterlab
[I 01:17:20.550 NotebookApp] JupyterLab application directory is /home/shared/anaconda3/share/jupyter/lab
[I 01:17:20.555 NotebookApp] Serving notebooks from local directory: /home/shared
[I 01:17:20.555 NotebookApp] 0 active kernels
[I 01:17:20.555 NotebookApp] The Jupyter Notebook is running at:
[I 01:17:20.555 NotebookApp] http://[all ip addresses on your system]:7000/?token=aad408a5bcc56f8a7d79db4e144507537e4cf927bd1ab6bc
[I 01:17:20.555 NotebookApp] Use Control-C to stop this server and shut down all kernels (twice to skip confirmation).
[C 01:17:20.556 NotebookApp]

Copy/paste this URL into your browser when you connect for the first time,
to login with a token:
http://localhost:7000/?token=aad408a5bcc56f8a7d79db4e144507537e4cf927bd1ab6bc
```

The important line (underlined in red) has the token for you to login from laptop. Replace the "localhost" part with your external IP address created in prior steps. In our example, the URL should be

```
http://35.185.240.182:7000/?token=aad408a5bcc56f8a7d79db4e144507537e4cf927bd1ab6bc
```

If there is no token, simply go to <http://35.185.240.182:7000>.

If you visit the above URL on your local browser, you should see something like the screen below.



The screenshot shows a Jupyter Notebook interface. At the top, there are tabs for 'Files', 'Running', and 'Clusters'. On the right, there is a 'Logout' button. Below the tabs, there is a message: 'Select items to perform actions on them.' A file list table follows:

		Name	Last Modified
<input type="checkbox"/>	<input type="checkbox"/>	/	
<input type="checkbox"/>	<input type="checkbox"/>	anaconda3	a day ago
<input type="checkbox"/>	<input type="checkbox"/>	DeepUbuntu	3 hours ago
<input type="checkbox"/>	<input type="checkbox"/>	setup.sh	7 hours ago

At the bottom of the table, there are buttons for 'Upload', 'New', and a refresh icon.

Submission: Transferring Files From Your Instance To Your Computer

When you are done with your assignments, run the submission script in your assignment folder to make a zip file. Please refer to specific instructions for each assignment.

Once you create the zip file, e.g. `assignment1.zip`, you will transfer the file from GCE instance to your local laptop. There is an [easy command](#) for this purpose:

```
gcloud compute scp <user>@<instance-name>:/path/to/assignment1.zip /local/path
```

For example, to download files from our instance to the current folder:

```
gcloud compute scp tonystark@cs231:/home/shared/assignment1.zip .
```

The transfer works in both directions. To upload a file to GCE:

```
gcloud compute scp /my/local/file tonystark@cs231:/home/shared/
```

Another (perhaps easier) option proposed by a student is to directly download the zip file from Jupyter. After running the submission script and creating `assignment1.zip`, you can download that file directly from Jupyter. To do this, go to Jupyter Notebook and click on the zip file, which will be downloaded to your local computer.

BIG REMINDER: Make sure you stop your instances!

Don't forget to stop your instance when you are done (by clicking on the stop button at the top of the page showing your instances). You can restart your instance and the downloaded software will still be available.

We have seen students who left their instances running for many days and ran out of credits. You will be charged per hour when your instance is running. This includes code development time. We encourage you to read up on Google Cloud, regularly keep track of your credits and not solely rely on our tutorials.